

Predictive Models of Methane Leak Rate

January 8, 2023

1 ENERES 131 Connector Assistant Research Project

Author: Richard Wang

1.1 Motivation

This research is originally conducted to see if the course staffs may improve one of the homework in the course. The original homework only has data from **CalEnviroScreen**. We would like to add another dataset and merge it with the CalEnviroScreen dataset, and use a feature from our new dataset as the target variable.

In the end, this new iteration of the homework was not implemented in the course, but I continued on the project. My goal in this research is to **find the best set of features that predicts methane leakage rate, so in the future when people try to prevent methane leakage, they can use these set of features as a starting point.**

1.2 Background on the new dataset - Methane Data

The data is provided by the **Environmental Defense Fund**. The portion we will be using contains only methane leakages collected in **Los Angeles**. - The measurements are the **methane leak rates** in a specific time span. - The measurements are measured in levels: *Low, Medium, High*. The feature is called **Leak_Cat**. - The measurements are measured **Per Geo Coordinate**, meaning every measurement is measured in a different geo coordinate - The time span is indicated in features **FIRST_D** and **LAST_D**

1.3 Research Question

Which features provide the most most predictive power for methane leakage rate?

1.3.1 Import Libraries

```
[ ]: # Run this cell to import the packages we will need
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import geopandas as gpd
```

```

import requests
import seaborn as sns
import xgboost as xgb

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.svm import SVR
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import GridSearchCV
%matplotlib inline

```

1.4 Part 1: EDA and Data Cleaning

1.4.1 Datasets used

- **LA_Shp**: Methane leak rate data provided by the EDF
- **ces20**: CalEnviroScreen environmental data

```
[2]: LA_shp = gpd.read_file('LosAngeles/Los_Angeles.shp')
ces20 = pd.read_csv('CES20/ces20.csv')
```

```
[3]: LA_shp.head()
```

```
[3]:
```

	FID_	FIRST_D	LAST_D	Lat	Lon	Leak_Cat	\
0	0.0	20150206	20150416	33.778466	-117.859590	Low	
1	1.0	20150322	20150406	33.780237	-117.863569	Low	
2	2.0	20150314	20150415	33.780511	-117.871238	Low	
3	3.0	20150211	20150406	33.780921	-117.841545	Low	
4	4.0	20150313	20150314	33.781181	-117.867509	Low	

```

                                geometry
0  POINT (-117.85959 33.77847)
1  POINT (-117.86357 33.78024)
2  POINT (-117.87124 33.78051)
3  POINT (-117.84154 33.78092)
4  POINT (-117.86751 33.78118)

```

```
[4]: ces20.head()
```

```
[4]:
```

	Census Tract	Total Population	California County	ZIP	City	\
0	6019001100	3174	Fresno	93706	Fresno	
1	6019000300	3609	Fresno	93706	Fresno	

2	6019000200	3167	Fresno	93706	Fresno
3	6019001500	2206	Fresno	93725	Fresno
4	6019000600	6161	Fresno	93721	Fresno

	Longitude	Latitude	Click for interactive map	CES 2.0 Score	\
0	-119.781696	36.709695	Click for map	89.22	
1	-119.801035	36.726462	Click for map	83.71	
2	-119.805504	36.735491	Click for map	83.47	
3	-119.717843	36.681600	Click for map	83.08	
4	-119.793357	36.743063	Click for map	82.95	

	CES 2.0 \nPercentile Range	...	Education Pctl	Linguistic Isolation	\
0	96-100% (highest scores)	...	95.60	21.6	
1	96-100% (highest scores)	...	91.05	18.3	
2	96-100% (highest scores)	...	93.95	16.2	
3	96-100% (highest scores)	...	90.48	19.5	
4	96-100% (highest scores)	...	90.09	16.9	

	Linguistic Isolation Pctl	Poverty	Poverty Pctl	Unemployment	\
0	83.66	77.500865	97.78	19.30	
1	78.34	81.204032	98.93	NaN	
2	74.04	86.828423	99.66	25.27	
3	80.68	62.746088	88.32	18.30	
4	75.49	88.680993	99.79	26.69	

	Unemployment Pctl	Pop. Char.	Pop. Char. Score	Pop. Char. Pctl
0	92.05	90.072268	9.360658	99.60
1	NaN	92.323243	9.594592	99.90
2	98.14	91.499039	9.508958	99.80
3	89.60	83.474281	8.674967	97.74
4	98.61	92.246260	9.586590	99.89

[5 rows x 55 columns]

1.4.2 Census Function

Although both datasets has **geo-coordinates** as features, merging on them results in an empty dataset because geo-coordinates are too specific, resulting in no matches between the two datasets.

So I decided to convert geo-coordinates to their belonging census tracts, and **merge the datasets on matching census tracts**. I did so using an API provided by the **Federal Communications Commission**

```
[5]: #Function to convert geo-coordinates to Census Tract
def Census(lat, lon):
    url = "https://geo.fcc.gov/api/census/area"
    r = requests.get(url, params = {'lat': lat, 'lon': lon})
    census = int(r.json()['results'][0]['block_fips'][1:11])
```

```
return census
```

1.4.3 Data cleaning and merge the two datasets

```
[6]: #Clean out nulls
LA_shp = LA_shp[LA_shp['Lat'] != 0]

#Add Census Tract column to LA_shp
LA_shp['Census Tract'] = LA_shp.apply(lambda df: Census(df.Lat, df.Lon), axis = 1)

#Merged data: CalEnviroScreen merged with LA Methane Leak on Census Tract
merged = pd.merge(ces20, LA_shp, how = 'inner', left_on = 'Census Tract',
                  right_on = 'Census Tract')

#Drop unused features and duplicates
new_merged = merged.drop(['FID_', 'FIRST_D', 'LAST_D', 'Lat', 'Lon',
                          'geometry'], axis = 1)
new_merged = new_merged.drop_duplicates()
```

Because every census tract may have multiple methane leak rate measurements, I grouped by census tracts, and used different aggregate functions to create the target variable: - **Leak_Rate_Average**: the average of leak rates in the census tract. Low = 0, Medium = 1, High = 2

```
[ ]: #group by census tract and creating new target variables
census_leak = new_merged[['Census Tract', 'Leak_Cat']]

census_leak['Leak_Rate'] = census_leak['Leak_Cat'].map({'Low': 0, 'Medium': 1,
              'High': 2})

census_leak = census_leak.groupby('Census Tract').agg({
    'Leak_Cat': lambda x: list(x),
    'Leak_Rate': 'mean'
})

census_leak = census_leak.rename(columns = {'Leak_Rate': 'Leak_Rate_Average'})
census_leak['Leakages'] = census_leak['Leak_Cat']
new_merged = new_merged.drop(['Click for interactive map', 'Hyperlink'], axis = 1)
```

Merge the new target variables with the cleaned CalEnviroScreen data.

```
[8]: new_merged = pd.merge(new_merged, census_leak, how = 'left', left_on = 'Census_
    Tract', right_on = 'Census Tract')
new_merged = new_merged.dropna()
```

1.5 Part 2: Modeling

1.5.1 Feature selection

For feature selection, I found the **linear correlation** between every feature and our target variable, and selected the features that has a **correlation > 0.1 or < -0.1**. After dropping the duplicating features, I am left with **14** features for the target variable.

```
[9]: new_merged = new_merged.drop(columns = ['Census Tract', 'California County',  
      ↪ 'ZIP'])  
corr = new_merged.corr()  
  
[10]: LeakAverageCorr = corr['Leak_Rate_Average']  
  
LeakAverageFeatures = LeakAverageCorr[(LeakAverageCorr > 0.1) |  
      ↪ (LeakAverageCorr < -0.1)]  
LeakAverageFeatures = LeakAverageFeatures.drop(labels = ['Diesel PM', 'Imp.  
      ↪ Water Bodies Pctl', 'Asthma Pctl', 'Low Birth Weight Pctl', 'Poverty Pctl',  
      ↪ 'Unemployment Pctl', 'Pop. Char. ', 'Pop. Char. Pctl', 'Leak_Rate_Average'])  
  
[11]: '''new_merged = new_merged.dropna()  
X = new_merged.loc[:, (new_merged.columns != 'Has_Leakage') &  
      (new_merged.columns != 'California County') &  
      (new_merged.columns != 'City') &  
      (new_merged.columns != 'Highest_Leakage') &  
      (new_merged.columns != 'Leakages') &  
      (new_merged.columns != 'Num_Leakages') &  
      (new_merged.columns != 'CES 2.0 \nPercentile Range')]  
X = X.set_index('Census Tract')'''  
X = new_merged[LeakAverageFeatures.index]  
y = new_merged['Leak_Rate_Average']  
y = y.astype('float')  
  
[12]: #train test split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,  
      ↪ random_state=42)
```

For modeling, I used **three** models: **Random Forest**, **XGBoost**, and a **Multi-layer Perceptron Neural Net**. For each model, I used **GridSearchCV** from sklearn for hyperparameter tuning and cross validation. I'm doing a **5-fold** CV.

1.5.2 Random Forest

```
[13]: forest = RandomForestRegressor(random_state = 42)  
  
param_grid = {
```

```

    'n_estimators': [5, 20, 50, 100],
    'max_depth': [5, 10, 15],
    'min_samples_split': [2, 6, 10],
    'min_samples_leaf': [1, 3, 4],
    'bootstrap': [True, False]
}

randomForest = GridSearchCV(estimator = forest, param_grid = param_grid, cv = 5,
    ↪ verbose = False)
randomForest.fit(X_train, y_train)
randomForestPred = randomForest.best_estimator_.predict(X_test)

```

1.5.3 XGBoost

```

[14]: xgb_model = xgb.XGBRFRegressor(random_state = 42)

param_grid = {
    'learning_rate': [ 0.01, 0.05],
    'max_depth': [5, 10, 15],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': np.arange(0.5, 1.0, 0.1),
    'n_estimators': [50, 100, 150],
}

xgboost = GridSearchCV(estimator = xgb_model, param_grid = param_grid, cv = 5,
    ↪ verbose = False)
xgboost.fit(X_train, y_train)
xgboostPred = xgboost.best_estimator_.predict(X_test)

```

1.5.4 Neural network - Multi-layer Perceptron

```

[ ]: mlp = MLPRegressor(random_state = 42)

param_grid = {
    'alpha': [0.0001, 0.001, 0.1],
    'activation': ['relu', 'tanh', 'logistic'],
    'hidden_layer_sizes': [(50,50,50), (50,100,50), (100,1)]
}

mlpModel = GridSearchCV(estimator = mlp, param_grid = param_grid, cv = 5,
    ↪ verbose = False)
mlpModel.fit(X_train, y_train)
mlpPred = mlpModel.best_estimator_.predict(X_test)

```

1.6 Part 3: Result

I used **three** methods for measuring the error for the model predictions: **Mean Squared Error**, **Mean Absolute Error**, and **R2 score**.

```
[16]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
#mse
randomForestMse = mean_squared_error(y_test, randomForestPred)
xgbMse = mean_squared_error(y_test, xgboostPred)
mlpMse = mean_squared_error(y_test, mlpPred)

#mae
randomForestMae = mean_absolute_error(y_test, randomForestPred)
xgbMae = mean_absolute_error(y_test, xgboostPred)
mlpMae = mean_absolute_error(y_test, mlpPred)

#r2
randomForestR2 = r2_score(y_test, randomForestPred)
xgbR2 = r2_score(y_test, xgboostPred)
mlpR2 = r2_score(y_test, mlpPred)
```

```
[17]: print('Random Forest MSE: ', randomForestMse)
print('XGBoost MSE: ', xgbMse)
print('MLP MSE: ', mlpMse)
print()
print('Random Forest MAE: ', randomForestMae)
print('XGBoost MAE: ', xgbMae)
print('MLP MAE: ', mlpMae)
print()
print('Random Forest R2: ', randomForestR2)
print('XGBoost R2: ', xgbR2)
print('MLP R2: ', mlpR2)
```

```
Random Forest MSE:  0.08370075707959367
XGBoost MSE:  0.10499091561245155
MLP MSE:  0.11442918253792207
```

```
Random Forest MAE:  0.25114127492488836
XGBoost MAE:  0.2398215596164976
MLP MAE:  0.26718253781113593
```

```
Random Forest R2:  0.2626809713438043
XGBoost R2:  0.0751362040431236
MLP R2:  -0.008005383255403231
```

2 Conclusion

This research project's main purpose is to first: **see if predicting methane leak rate with environmental data from CalEnviroScreen is possible**, and if so, **what are the best features and models for it?** With the results from the three models I've tested, the three models performed similarly in predicting methane leak rates.

2.0.1 Most Predictive Features

These are the 14 features and their correlation with leakage rate, ranked from **most predictive to least**: - Unemployment: **.268081** - Pop. Char. Score: **0.228498** - Asthma: **0.211940** - Poverty: **0.187427** - Imp. Water Bodies: **-0.185010** - CES 2.0 Score: **0.173223** - Low Birth Weight: **0.169838** - Education Pctl: **0.145919** - Traffic Pctl: **-0.144668** - Age Pctl: **-0.128569** - Pesticides Pctl: **-0.122472** - Cleanup Sites Pctl: **0.121789** - Linguistic Isolation Pctl: **0.120849** - Diesel PM Pctl: **0.103990**

2.0.2 Potential issues and flaws

- **Size of data:** the biggest issue is I am working with a dataset that has around 60 measurements due to the mismatches in the merging process, so the results is potentially very biased
- **Nature of the datasets:** the two datasets are measured different, one measured **per census**, one measured **per coordinate**. Although the API resolved the issue, it largely reduced the number of measurements I can work with

2.0.3 Future improvements

Add data: try a dataset that is measured in **per census**, and also cover more California states other than Los Angeles. A good place to start is to find data on the best features for predicting leak rates, as listed above. The data with those features that is also measured in per census will work the best.